

```
In [1]: import pandas as pd

In [2]: adv_df = pd.read_csv('Advertising.csv', index_col=0)

In [3]: adv_df.head()

Out[3]:
   TV  radio  newspaper  sales
1 230.1   37.8     69.2   22.1
2  44.5   39.3     45.1   10.4
3  17.2   45.9     69.3    9.3
4 151.5   41.3     58.5   18.5
5 180.8   10.8     58.4   12.9

In [8]: adv_df.size

Out[8]: 800

In [9]: adv_df.shape

Out[9]: (200, 4)

In [10]: adv_df.columns

Out[10]: Index(['TV', 'radio', 'newspaper', 'sales'], dtype='object')

In [11]: x_features=adv_df[['newspaper', 'radio', 'TV']]

In [12]: x_features.head()

Out[12]:
   newspaper  radio    TV
1     69.2   37.8  230.1
2     45.1   39.3   44.5
3     69.3   45.9   17.2
4     58.5   41.3  151.5
5     58.4   10.8  180.8

In [13]: y_target=adv_df[['sales']]

In [14]: y_target.head()

Out[14]:
   sales
1    22.1
2    10.4
3     9.3
4    18.5
5    12.9

In [16]: x_features.shape

Out[16]: (200, 3)

In [17]: y_target.shape

Out[17]: (200, 1)

In [18]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_features, y_target, random_state=1)

In [19]: print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

(150, 3)
(150, 1)
(50, 3)
(50, 1)

In [20]: from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(x_train,y_train)

Out[20]: LinearRegression()

In [22]: print(linreg.intercept_)
print(linreg.coef_)

[2.87696662]
[[0.00345046  0.17915812  0.04656457]]

In [23]: y_pred=linreg.predict(x_test)
y_pred

Out[23]: array([[21.70910292],
 [16.41055243],
 [ 7.60955058],
 [17.80769552],
 [18.6146359 ],
 [23.83573998],
 [16.32488681],
 [13.43225536],
 [ 9.17173403],
 [17.333853  ],
 [14.44479482],
 [ 9.83511973],
 [17.18797614],
 [16.73086831],
 [15.05529391],
 [15.61434433],
 [12.42541574],
 [17.17716376],
 [11.08827566],
 [18.00537501],
 [ 9.28438889],
 [12.98458458],
 [ 8.79950614],
 [10.42382499],
 [11.3846456 ],
 [14.98082512],
 [ 9.78853268],
 [19.39643187],
 [18.18099936],
 [17.12807566],
 [21.54670213],
 [14.69809481],
 [16.24641438],
 [12.32114579],
 [19.92422501],
 [15.32498602],
 [13.88726522],
 [10.03162255],
 [20.93105915],
 [ 7.44936831],
 [ 3.64695761],
 [ 7.22020178],
 [ 5.9962782 ],
 [18.43381853],
 [ 8.39408045],
 [14.08371047],
 [15.02195699],
 [20.35836418],
 [20.57036347],
 [19.60636679]])

In [25]: from sklearn import metrics
import numpy as np

In [26]: print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))

1.404651423032897

In [29]: print('True',y_test.values[0:10])
print()
print('Pred',y_pred[0:10])

True [[23.8]
 [16.6]
 [ 9.5]
 [14.8]
 [17.6]
 [25.5]
 [16.9]
 [12.9]
 [10.5]
 [17.1]]

Pred [[21.70910292]
 [16.41055243]
 [ 7.60955058]
 [17.80769552]
 [18.6146359 ]
 [23.83573998]
 [16.32488681]
 [13.43225536]
 [ 9.17173403]
 [17.333853  ]]
```

```
In [ ]:
```